

Installing CDH (Cloudera Distribution Including Apache Hadoop) 5.5 on Mac OSX

1. Dependencies

- a. Install the Java version that is supported for the CDH version you are installing.
- b. Enable ssh on your mac by turning on remote login. You can find this option under your toolbar's System Preferences > Sharing.
 1. Check the box for Remote Login to enable the service.
 2. Allow access for: "Only these users: Administrators"
- c. Enable password-less ssh login to localhost for MRv1 and HBase.
 1. Generate an rsa or dsa key.
 1. `ssh-keygen -t rsa -P ""`
 2. Continue through the key generator prompts (use default options).
 2. Test: `ssh localhost`

2. CDH

The CDH tarballs are very nicely packaged and easily downloadable from Cloudera's repository. Download and explode the tarballs in a lib directory where you can manage latest versions with a simple symlink as the following. Although Mac OSX's "Make Alias" feature is bi-directional, do not use it, but instead use your command-line `ln -s` command, such as `ln -s source_file target_file`.

- `/Users/najus/cloudera/`
- `cdh5.1/`
- `hadoop -> /Users/najus/cloudera/lib/hadoop-2.3.0-cdh5.1.0`
- `hbase -> /Users/najus/cloudera/lib/hbase-0.98.1-cdh5.1.0`
- `hive -> /Users/najus/cloudera/lib/hive-0.12.0-cdh5.1.0`
- `zookeeper -> /Users/najus/cloudera/lib/zookeeper-3.4.5-cdh4.7.0`
- `ops/`
- `dn`
- `logs/hadoop, logs/hbase, logs/yarn`
- `nn/`
- `pids`
- `tmp/`
- `zk/`

Now, set your environment properties according to the paths where you've exploded your tarballs.

```
CDH="cdh5.1"
export HADOOP_HOME="/Users/najus/cloudera/${CDH}/hadoop"
export HBASE_HOME="/Users/najus/cloudera/${CDH}/hbase"
export HIVE_HOME="/Users/najus/cloudera/${CDH}/hive"
export HCAT_HOME="/Users/najus/cloudera/${CDH}/hive/hcatalog"

export
PATH=${JAVA_HOME}/bin:${HADOOP_HOME}/bin:${HADOOP_HOME}/sbin:${ZK_HOME}/bin:${HBASE_HOME}/bin:${HIVE_HOME}/bin:${HCAT_HOME}/bin:${M2_HOME}/bin:${ANT_HOME}/bin:${PATH}
```

Update your main Hadoop configuration files, as shown in the sample files below.

```
1. $HADOOP_HOME/etc/hadoop/core-site.xml
<configuration>
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:8020</value>
  <description>The name of the default file system. A URI whose
    scheme and authority determine the FileSystem implementation. The
    uri's scheme determines the config property (fs.SCHEME.impl) naming
    the FileSystem implementation class. The uri's authority is used to
    determine the host, port, etc. for a filesystem.</description>
</property>
<property>
  <name>hadoop.tmp.dir</name>
  <value>/Users/najus/cloudera/ops/tmp/hadoop-${user.name}</value>
  <description>A base for other temporary directories.</description>
</property>
<property>
  <name>io.compression.codecs</name>

  <value>org.apache.hadoop.io.compress.GzipCodec,org.apache.hadoop.io.compress.DefaultCodec,org.apache.hadoop.io.compress.BZip2Codec,org.apache.hadoop.io.compress.SnappyCodec</value>

  <description>A comma-separated list of the compression codec classes that can
    be used for compression/decompression. In addition to any classes specified
    with this property (which take precedence), codec classes on the classpath
    are discovered using a Java ServiceLoader.</description>
</property>
</configuration>
```

2. \$HADOOP_HOME/etc/hadoop/hdfs-site.xml

```
<configuration>
<property>
  <name>dfs.namenode.name.dir</name>
  <value>/Users/najus/cloudera/ops/nn</value>
  <description>Determines where on the local filesystem the DFS name node
    should store the name table(fsimage). If this is a comma-delimited list
    of directories then the name table is replicated in all of the
    directories, for redundancy. </description>
</property>
<property>
  <name>dfs.datanode.data.dir</name>
  <value>/Users/najus/cloudera/ops/dn</value>
  <description>Determines where on the local filesystem an DFS data node
    should store its blocks. If this is a comma-delimited
    list of directories, then data will be stored in all named
    directories, typically on different devices.
    Directories that do not exist are ignored.
  </description>
</property>
<property>
  <name>dfs.datanode.http.address</name>
  <value>localhost:50075</value>
  <description>
    The datanode http server address and port.
    If the port is 0 then the server will start on a free port.
  </description>
</property>
<property>
  <name>dfs.replication</name>
  <value>1</value>
  <description>Default block replication.
    The actual number of replications can be specified when the file is
    created.
    The default is used if replication is not specified in create time.
  </description>
</property>
</configuration>
```

3. \$HADOOP_HOME/etc/hadoop/yarn-site.xml

```
<configuration>
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
  <description>the valid service name should only contain a-zA-Z0-9_ and can
    not start with numbers</description>
</property>
</configuration>
```

```
<name>yarn.log-aggregation-enable</name>
<value>true</value>
<description>Whether to enable log aggregation</description>
</property>
<property>
  <name>yarn.nodemanager.remote-app-log-dir</name>
  <value>hdfs://localhost:8020/tmp/yarn-logs</value>
  <description>Where to aggregate logs to.</description>
</property>
<property>
  <name>yarn.nodemanager.resource.memory-mb</name>
  <value>8192</value>
  <description>Amount of physical memory, in MB, that can be allocated
    for containers.</description>
</property>
<property>
  <name>yarn.nodemanager.resource.cpu-vcores</name>
  <value>4</value>
  <description>Number of CPU cores that can be allocated
    for containers.</description>
</property>
<property>
  <name>yarn.scheduler.minimum-allocation-mb</name>
  <value>1024</value>
  <description>The minimum allocation for every container request at the RM,
    in MBs. Memory requests lower than this won't take effect,
    and the specified value will get allocated at minimum.</description>
</property>
<property>
  <name>yarn.scheduler.maximum-allocation-mb</name>
  <value>2048</value>
  <description>The maximum allocation for every container request at the RM,
    in MBs. Memory requests higher than this won't take effect,
    and will get capped to this value.</description>
</property>
<property>
  <name>yarn.scheduler.minimum-allocation-vcores</name>
  <value>1</value>
  <description>The minimum allocation for every container request at the RM,
    in terms of virtual CPU cores. Requests lower than this won't take effect,
    and the specified value will get allocated the minimum.</description>
</property>
<property>
  <name>yarn.scheduler.maximum-allocation-vcores</name>
  <value>2</value>
  <description>The maximum allocation for every container request at the RM,
    in terms of virtual CPU cores. Requests higher than this won't take effect,
    and will get capped to this value.</description>
</property>
```

```
</property>
</configuration>
```

4. \$HADOOP_HOME/etc/hadoop/mapred-site.xml

```
<configuration>
<property>
  <name>mapreduce.jobtracker.address</name>
  <value>localhost:8021</value>
</property>
<property>
  <name>mapreduce.jobhistory.done-dir</name>
  <value>/tmp/job-history</value>
  <description></description>
</property>
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
  <description>The runtime framework for executing MapReduce jobs.
  Can be one of local, classic or yarn.
  </description>
</property>
<property>
  <name>mapreduce.map.cpu.vcores</name>
  <value>1</value>
  <description>
    The number of virtual cores required for each map task.
  </description>
</property>
<property>
  <name>mapreduce.reduce.cpu.vcores</name>
  <value>1</value>
  <description>
    The number of virtual cores required for each reduce task.
  </description>
</property>
<property>
  <name>mapreduce.map.memory.mb</name>
  <value>1024</value>
  <description>Larger resource limit for maps.</description>
</property>
<property>
  <name>mapreduce.reduce.memory.mb</name>
  <value>1024</value>
  <description>Larger resource limit for reduces.</description>
</property>
<property>
  <name>mapreduce.map.java.opts</name>
```

```

    <value>-Xmx768m</value>
    <description>Heap-size for child jvms of maps.</description>
</property>
<property>
    <name>mapreduce.reduce.java.opts</name>
    <value>-Xmx768m</value>
    <description>Heap-size for child jvms of reduces.</description>
</property>
<property>
    <name>yarn.app.mapreduce.am.resource.mb</name>
    <value>1024</value>
    <description>The amount of memory the MR AppMaster needs.</description>
</property>
</configuration>

```

5. \$HADOOP_HOME/etc/hadoop/hadoop-env.sh (indicated properties only)

```

# Where log files are stored. $HADOOP_HOME/logs by default.
export HADOOP_LOG_DIR="/Users/najus/cloudera/ops/logs/hadoop"
export YARN_LOG_DIR="/Users/najus/cloudera/ops/logs/yarn"

# The directory where pid files are stored when processes run as daemons. /tmp
by default.
export HADOOP_PID_DIR="/Users/najus/cloudera/ops/pids"
export YARN_PID_DIR=${HADOOP_PID_DIR}

```

6. \$HBASE_HOME/conf/hbase-site.xml

```

<configuration>
<property>
    <name>hbase.cluster.distributed</name>
    <value>true</value>
    <description>The mode the cluster will be in. Possible values are
        false for standalone mode and true for distributed mode. If
        false, startup will run all HBase and ZooKeeper daemons together
        in the one JVM.
    </description>
</property>
<property>
    <name>hbase.tmp.dir</name>
    <value>/Users/najus/cloudera/ops/tmp/hbase-${user.name}</value>
    <description>Temporary directory on the local filesystem.
        Change this setting to point to a location more permanent
        than '/tmp' (The '/tmp' directory is often cleared on
        machine restart).
    </description>
</property>
</property>

```

```

    <name>hbase.zookeeper.property.dataDir</name>
    <value>/Users/najus/cloudera/ops/zk</value>
    <description>Property from ZooKeeper's config zoo.cfg.
    The directory where the snapshot is stored.
    </description>
  </property>
  <property>
    <name>hbase.rootdir</name>
    <value>hdfs://localhost:8020/hbase</value>
    <description>The directory shared by region servers and into
    which HBase persists. The URL should be 'fully-qualified'
    to include the filesystem scheme. For example, to specify the
    HDFS directory '/hbase' where the HDFS instance's namenode is
    running at namenode.example.org on port 9000, set this value to:
    hdfs://namenode.example.org:9000/hbase. By default HBase writes
    into /tmp. Change this configuration else all data will be lost
    on machine restart.
    </description>
  </property>
</configuration>

```

```

7. $HBASE_HOME/conf/hbase-env.sh
# Where log files are stored. $HBASE_HOME/logs by default.
# Where log files are stored. $HBASE_HOME/logs by default.
export HBASE_LOG_DIR="/Users/najus/cloudera/ops/logs/hbase"

# The directory where pid files are stored. /tmp by default.
export HBASE_PID_DIR="/Users/najus/cloudera/ops/pids"

# Tell HBase whether it should manage its own instance of Zookeeper or not.
export HBASE_MANAGES_ZK=true

```

3. Putting it all together

By now, we should have accomplished setting up HDFS, YARN, and HBase. These are the bare essentials for getting your local machine ready for running MapReduce jobs and building applications on HBase. In the next few steps, we will start/stop the services and provide examples to ensure each service is operating correctly. The steps are listed in the specific order for initialization in order to adhere to dependencies. The order could be reversed for halting the services.

Service HDFSs

NameNode

format: hdfs namenode -format

start: hdfs namenode

stop: Ctrl-C

url: <http://localhost:50070/dfshealth.html>

DataNode

start: hdfs datanode

stop: Ctrl-C

url: <http://localhost:50075/browseDirectory.jsp?dir=%2F&nnaddr=127.0.0.1:8020>

Test

hadoop fs -mkdir /tmp

hadoop fs -put /path/to/local/file.txt /tmp/

hadoop fs -cat /tmp/file.txt

Service YARN

ResourceManager

start: yarn resourcemanager

stop: Ctrl-C

url: <http://localhost:8088/cluster>

NodeManager

start: yarn nodemanager

stop: Ctrl-C

url: <http://localhost:8042/node>

4. Running the program on Hadoop

In terminal, we use “hadoop jar” command to run our jar on the hadoop system.

For this, we create an “input” folder using `hdfs dfs -mkdir input` command where we put our input data file.

To run the jar file we do the following:

```
hadoop jar /Users/najus/MUM/Big\  
Data/workspace/Crystalball/target/crystalball-0.0.1-SNAPSHOT.jar  
edu.mum.wordcount.WordCount /input/sample.txt /output
```